AIPS Mini Hackathon Judging Rubric

Case: Automating the Learning Support Tutor's Workday I AI Prototyping Studio, University of Melbourne

How to score

Each team is scored out of 100 across seven weighted criteria. Feasibility and prototype quality carry the most weight, because the brief asks for tasks that can be partially automated in a real tutoring centre, not a full replacement of the tutor.

For each criterion, judges pick a level from 1 to 5 using the descriptors below. Scores of 2 sit between the Weak and Fair descriptors. Convert to points with: **points = level × (weight ÷ 5)**. So a level 4 on Feasibility (weight 20) gives 16 points.

If two teams tie, the higher Feasibility score wins. If still tied, the higher Impact score wins. Before judging starts, all judges should score one sample pitch together so that everyone agrees on what a 3 looks like.

Rubric

Criterion	Weight	What it assesses	Excellent (5)	Good (4)	Fair (3)	Weak (1)
1. Problem selection & insight	**15**	Whether the team understood the tutor's workday and chose a task where AI genuinely helps.	Targets specific numbered steps from the brief and shows why they are a bottleneck, e.g. the same marking result entered into all three exhibit sheets in step 8. Explains what the tutor still has to do.	Clear link to one or two steps in the workflow. Reasoning for the choice is sound but light on detail.	Problem is relevant to tutoring but only loosely tied to the brief's steps.	Generic idea ("an AI tutor") with no reference to the actual workflow.
2. Creativity	**15**	Originality of the approach, and whether it uses the centre's existing data in a smart way.	A fresh angle judges haven't seen from other teams. Makes clever use of what's already recorded, such as using the W vs C marks to separate knowledge gaps from careless errors.	Some original thinking in at least one part of the solution.	A familiar idea done competently.	Essentially a chatbot wrapper with nothing specific to this case.
3. Feasibility	**20**	Could the centre actually adopt this soon, with its current tools, budget and staff?	Works with tools already in use (Google Sheets, Drive, Classroom, Work WeChat). Realistic costs. A tutor could pilot it within weeks. Main technical risks named, with a fallback.	Mostly fits current tools. One or two gaps the team acknowledges.	Plausible, but needs a major change in tools or habits that the team doesn't address.	Depends on things that don't exist yet, or would require replacing the centre's systems.
4. Prototype quality	**20**	How much actually works in the demo.	Live, working demo on realistic data (e.g. data modelled on Exhibits 1 to 3). Handles at least one messy case, like a blank answer or a message in Mandarin.	Working demo of the core flow. Some parts are mocked but clearly labelled.	Partial demo. Key steps are slides or screenshots.	Concept only, or the demo fails.
5. Impact	**15**	How much time or quality the solution gives back to the tutor and students.	Quantified estimate of time saved against the 14:00 to 19:00 shift, with a believable basis. Also shows a benefit to students or parents.	Clear benefit with a rough estimate.	Benefit described but not measured.	Unclear who benefits or by how much.
6. Responsible AI	**10**	Safe, sensible use of AI with student data, given most students are minors.	Tutor keeps the judgement calls (passing a Demo, scheduling a Retest, messages to parents). Handles student data privacy and explains what happens when the AI is wrong.	Human review is built in. Privacy mentioned.	Some awareness of risks, but no design response.	AI makes decisions about students unchecked, or risks are ignored.
7. Pitch	**5**	Clarity of communication within the time limit.	Problem, solution and demo land clearly, on time. Strong answers to judges' questions.	Clear, minor timing or flow issues.	Understandable, but hard to follow in places.	Unclear, or well over time.
Total	**100**					